



Global Preferences for Hate Speech Moderation

Public registration ▼ Updates ▼



Metadata

Study Information



Hypotheses

Hypotheses on vignette attribute effects moderated by respondent characteristics

H1: Respondents who agree (vs. disagree) with the issue position of the target are more likely to perceive messages as being hate speech and offensive speech, and to prefer platform, legal, and employer action against the message/sender.

H2: Respondents who share gender or ethnicity with the target are more likely to perceive messages as being hate speech and offensive speech, and to prefer platform, legal, and employer action against the message/sender.

H3: Respondents who share gender or ethnicity with the sender are less likely to perceive messages as being hate speech and offensive speech, and to prefer platform, legal, and employer action against the message/sender.

Hypotheses on the framing experiment

H4: If the vignette evaluation is framed as a moderation task for a platform that promotes free speech and takes a critical stance against any form of censorship (vs. a neutral frame), respondents are less likely to prefer platform action against the message sender but are as likely to perceive messages as being hate speech and offensive speech, and to prefer legal action against the message/sender.

H5: If the vignette evaluation is framed as a moderation task for a platform that vows to protect vulnerable users from hate speech and takes an aggressive stance against hateful content on their pages (vs. a neutral frame), respondents are more likely to prefer platform action against the message sender but are as likely to perceive messages as being hate speech and offensive speech, and to prefer legal action against the message/sender.

Hypotheses on question order experiment

H6: If respondents are exposed to the vignette evaluation task before (vs. after) they get to answer

questions about governance of speech, they are more willing to have speech restricted, and see more responsibility for action against hate speech in the hands of online services, lawmakers, the justice system, and employers.

Design Plan

Study type

Experiment - A researcher randomly assigns treatments to study subjects, this includes field or lab experiments. This is also known as an intervention experiment and includes randomized controlled trials.

Blinding

For studies that involve human subjects, they will not know the treatment group to which they have been assigned.

Is there any additional blinding in this study?

No response

Study design

We conduct a vignette design embedded in a cross-sectional survey fielded in multiple countries (Brazil, Colombia, Germany, India, Indonesia, Nigeria, Philippines, Poland, Turkey, United Kingdom, United States). The rationale behind the selection of countries was to maximize coverage of different cultural contexts across regions of the world (Europe: 4, America: 3, Asia: 3, Africa: 1) as well as variation in free speech norms as reported in the Future of Free Speech Survey (Skaaning and Krishnarajan, 2021). The US survey served as the baseline for translations and contextual adaptations for all other countries. The translations were performed by native speakers who had familiarity with the respective country contexts.

We construct vignettes in forms of social media posts, mimicking actual cases of toxic conversations on social media to bolster external validity. The conversation vignettes are synthetically generated based on several attributes, including the topic, topic stance, speech type, target identity, and sender identity. Vignettes are constructed based on random assignment of attribute levels. The study design is informed by prior work of ours (Munzert et al., 2022) as well as more recent work (Rasmussen, 2022).

Respondents are then asked to state (1) whether they consider the message (a) hate speech, (b) offensive, but no hate speech, or (c) neither offensive nor hate speech, (2) whether they think the statement should be removed from the platform, (3) whether they think the sender should be banned from the platform, (4) whether they think the sender should face legal consequences, and (5) whether they think the sender should face consequences from their employer. Each respondent is asked to evaluate eight vignettes.

Furthermore, the vignette task is embedded in different, randomly assigned frames to test whether respondents can be framed to take a rather firm or soft stance towards hateful content.

Finally, the vignette task is coupled with a question order experiment to test for the consequences of hate speech exposure on preferences for regulation and governance of online hate speech.

Lin, Winston. 2013. "Agnostic Notes on Regression Adjustments to Experimental Data: Reexamining Freedman's Critique." *Annals of Applied Statistics* 7(1):295–318.

Munzert, S., Traunmüller, R., Guess, A, Barberà, P., Yang, J. (2022). What Should We Be Allowed to Post? Citizens' Preferences for Online Hate Speech Regulation. <https://osf.io/yvcdq>

Rasmussen, J. (2022). How the Public Form Judgements of Online Hate Speech. <https://osf.io/7xvq6/>

Skaaning, S. E., & Krishnarajan, S. (2021). Who Cares About Free Speech?: Findings from a Global Survey of Support for Free Speech. https://futurefreespeech.com/wp-content/uploads/2021/06/Report_Who-cares-about-free-speech_21052021.pdf

No files selected

Randomization

Vignette attribute levels are assigned randomly. Vignette decks are balanced across several attribute values so that each respondent gets to see a mix of vignettes that vary along message, sender, and target characteristics. Specific sender and target ethnicity characteristics are oversampled depending on the country context to account for prevalent ethnicities in the different settings. Details on vignette construction and sampling are documented in the vignette construction repository [vignette-construction.zip](#).

Sampling Plan

Existing Data

Registration prior to creation of data

Explanation of existing data

No response

Data collection procedures

Participants will be recruited using Facebook Ads. No incentives will be given. We will use targeting based on age (18-40 and 41-65+), sex (male, female), and education (below college, at least college education) where necessary to attain a sample that is balanced on the marginals across those categories. We show Facebook Ads both on the Facebook page of “The Digital Citizen Lab” and on the Instagram account “digitalcitizenlab”. We select “highest volume” as campaign bid strategy and use Advantage+ placements.

We use two different ad versions. The neutral ad consisting of the text “Do you want to participate in our survey? Hertie School researchers want to hear your opinions.” and a neutral image of a pencil, a question mark or a microphone. A second ad consists of the text “What should people be allowed to say on social media? Hertie School researchers want to hear your opinions.” and an image of a stop sign, a stop signaling hand in front of face, or a phone with social media apps. We may change the images within an ad over the course of the recruitment period to avoid visual fatigue among potential participants as recommended by Meta. The audiences consist of people living in the respective country, speaking the language in which the survey and ad has been translated.

Ads will be delivered through the same campaign, in which we will use conversion as the campaign objective. Data will be collected from mid-November to December 2022 until either the campaign

budget is exhausted or the project funding period has ended (end of December). Once participants click on the “learn more” action button in the advertisement, they will be directed to our survey, hosted on a Qualtrics page. Before starting the survey, participants will be informed about the length, objective and risks of the survey and they will be asked to provide consent. The participants are debriefed at the end of the survey.

No files selected

Sample size

We target a sample size of at least 1,500 respondents per country.

Sample size rationale

The sample size is constrained by research budget and will be guided by Facebook’s Ad conversion toolkit.

Stopping rule

Data will be collected from mid-November to December 2022 until either the campaign budget is exhausted or the project funding period has ended (end of December).

Variables

Manipulated variables

The vignettes are constructed in a way that mimics short conversations on a popular social media platform. The conversations are manipulated along a set of target, sender, and message characteristics. Those include: (1) topic of the conversation, (2) target’s gender, (3) target’s ethnicity, (4) target’s topic stance, (5) sender’s gender, (6) sender’s ethnicity, (7) sender’s message target, (8) sender’s message type, (9) sender’s message severity.

The messages have been generated by the researchers and adapted to the topic and country context by translators who were closely supervised during the process. Technical details of the vignette construction as well as the raw material is documented in the vignette construction repository [vignette-construction.zip](#).

Prior to exposing respondents to the set of hate speech vignettes, they are made aware of the nature of the content that they were about to see and encouraged to skip the task if they feel uncomfortable.

Respondents are assigned one of three frames of the evaluation task with equal probability ($\frac{1}{3}$). They read as follows:

Neutral frame: “You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content. How would you act if you had to decide what happens with the messages?”

Free speech promotion frame: “You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content. Imagine you were a content moderator for a social platform that vows to promote free speech and takes a critical stance against any form of censorship. How would you act if you had to decide what happens with the messages?”

User protection frame: “You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content. Imagine you were a content moderator for a social platform that vows to protect their users from hate speech and takes a committed stance against hateful content on their pages. How would you

act if you had to decide what happens with the messages?”

Finally, a battery of questions is presented to the respondents either before or after the vignettes evaluation task, with equal probability (½).

- Questionnaire-USA.pdf
(<https://osf.io/project/2ec8m/files/osfstorage/63749467b076650485617afd>)
- vignettes-construction.zip
(<https://osf.io/project/2ec8m/files/osfstorage/63749466b076650485617afb>)

Measured variables

Five outcome variables are measured for each of the vignettes: Perceived offensiveness/hate speech, preference for removal of message, preference for ban of sender, preference for legal action against sender, preference for employer action against sender.

Do you think that the flagged post is hate speech, offensive speech, or neither of them?

- Hate speech
- Offensive but not hate speech
- Neither offensive nor hate speech

Do you think the flagged post should be removed from the platform or not?

- Post should be removed
- Post should not be removed

Do you think the sender of the flagged post should be banned from the platform or not?

- Person should be banned
- Person should not be banned

Do you think that the sender of the flagged post should face legal consequences, such as a fine?

- Person should face legal consequences
- Person should not face legal consequences

Do you think that the sender should face consequences from their employer, such as being fired?

- Person should face consequences by employer
- Person should not face consequences by employer

In addition, several additional outcome measures are taken (either just before or after the vignettes evaluation task:

- Free speech trade-offs (three items)
- Group-specific responsibility to act against online hate speech (six groups)
- Government/companies regulation trade-off

After the vignette evaluation we provide a question that serves as a manipulation check. It asks respondents about the nature of the social platform they had in mind when answering evaluating the vignettes.

At the very end of the survey we pose an open question and ask respondents to provide their personal definition of “hate speech”. Furthermore, we ask respondents to put in their own words what - if anything - should not be allowed to say on social media.

Furthermore, several pre-treatment variables are measured:

- Gender (female/male/other)

- Year of birth; derived age and collapsed into categories 18-29, 30-49, 50-69, 70+
- Education; measured, as highest level of school completed (country-specific categories) and collapsed into three levels (low/intermediate/high)
- Country of residence (not via survey item, but survey taken)
- Visual self-identification with ethnicities; measured using a visual avatar selection task matching avatars used in the vignettes
- Self-reported minority status (yes/no)
- Political interest; measured with four levels and collapsed into two (low/high)
- Political partisanship; measured as country-specific self-identification with parties
- Political ideology; measured with 11-point left/liberal - right/conservative scale; collapsed into three levels (left/center/right)
- Agreement with target statements (tend to agree/tend to disagree/unsure)
- Freedom of speech perception now and then (general and personal), measured with three levels (more free/less free/about same)
- Experience with hate and offensive speech, five items
- Online hostile engagement, three items

Finally, a couple of indices are created on the basis of some survey items:

- We will measure respondents' preference to silence others by conducting a principal components analysis (PCA) of their responses to two topic-specific items and taking the first principal component.
- We will measure respondents' empathy by conducting a principal components analysis (PCA) of their responses to three topic-specific items and taking the first principal component.
- We will measure general experience with hate and offensive speech by computing the average score on five topic-specific items.
- We will measure online hostile engagement by computing the average score on three topic-specific items.

A copy of the full US questionnaire can be found in "Questionnaire-USA".

No files selected

Indices

No response

No files selected

Analysis Plan

Statistical models

General modeling strategy. To analyze our vignette survey experiment, we use hierarchical modeling with person, message, and country random effects (in pooled analyses). This allows us to deal with several key challenges (Gelman, 2006). First, by including varying intercepts for individual respondents, we accommodate the nesting of vignette judgments within individuals, producing meaningful standard errors. Second, random intercepts for messages help us account for any message-specific effects. Note that while the set of message attributes is identical across all countries, the set of vignettes is not because sender and target characteristics vary due to oversampling of particular ethnic groups by country. Third, partial pooling helps address the multiple comparisons problem.

For all vignette evaluation outcomes (perceived as hate speech vs not, perceived as offensive vs. not, post removed or not, sender banned or not, sender facing legal consequences or not, sender facing employer consequences or not), we will use linear model specifications. Using a linear model

specification for binary outcomes (action chosen = 1 vs. action not chosen = 0) yields a straightforward interpretation of coefficients in terms of probabilities of agreeing with the treatment of controversial or hate speech.

The response time measure on all vignette pages will be used to report, in addition to the models using the full sample, models that are run only using vignette evaluations for which respondents spent at least 5 seconds on the respective page.

Specification of descriptive models. In addition to models designed to test the pre-registered hypotheses (see below), we will present two sets of models that describe (1) the effects of vignette attribute characteristics and (2) the effects of respondent characteristics on evaluation outcomes. In (1), we will provide fully specified models that include all attribute characteristics except for sender/target names (sender/target identity is only measured using ethnicity + gender). In (2), we will provide fully specified models that include all pre-treatment measures except for treatment status. Results will be presented visually reporting marginal means as well as in tabular format.

Hypothesis tests. Hypotheses will be tested on the pooled sample. To test for the moderating effect of respondent characteristics on hate speech regulation preferences (H1-H3), we construct indicators of overlap based on respondents' responses on relevant pre-vignette items and vignette characteristics. Specifically, to test H1 (target issue position agreement), we create a respondent-vignette-specific indicator whether the respondent agrees or does not agree with the target position. To test H2 and H3, we compare the characteristics of the avatar the respondent identified with (gender, ethnicity) with the corresponding characteristics of the sender/target avatar, and again construct respondent-vignette-specific indicators (respondent-target identity overlap yes/no).

Hypotheses on the framing experiment (H4/H5) will be tested by regressing the evaluation outcomes (a) only on one of the treatment group indicators (free speech vs. neutral frame, protect users vs. neutral frame) and (b) on the treatment group indicator and a covariate-adjusted model using Lin's (2013) saturated regression approach including all pre-treatment covariates.

The manipulation check question will be used to report, in addition to a model using the full sample, a model that is run only using treated respondents who answered the manipulation check question in accordance with their treatment condition, and all respondents in the neutral condition irrespective of how they answered the manipulation check question. The response time measure on the vignette task frame will be used to report, in addition to a model using the full sample, a model that is run only using treated respondents who spent at least 10 seconds on the survey page providing the frame.

The hypothesis on the question order experiment (H6) will be tested by regressing the outcomes on the governance of speech preference questions (a) only on the pre/post treatment indicator (sparse model) and (b) on the pre/post treatment indicator and a covariate-adjusted model using Lin's (2013) saturated regression approach including all pre-treatment covariates.

No files selected

Transformations

See above "Measured variables".

Inference criteria

We will rely on classical frequentist statistical inference and the conventional $p < .05$ cut off for statistical significance. All tests will be two-tailed.

Data exclusion

We will exclude speeders from the analysis, whom we define as respondents who spent less than five minutes on the entire survey.

Missing data

We will rely on model specific listwise deletion of observations with missing variable values.

Exploratory analysis

In addition to our descriptive analyses as well as the analyses guided by our hypotheses, we will conduct a range of exploratory analyses and additional analyses to assess the robustness of results.

First, we will explore the variation of effects by country context, subsetting the data accordingly. Second, we will explore the variation of attribute effects by selected respondent characteristics, also by subsetting the data along selected variables. Third, we will explore how selected attributes moderate the effects of other selected attributes.

Furthermore, we will report a range of descriptive statistics on the measured outcomes, including a comparison of the relative prevalence of different evaluation outcomes (post removed, sender band, legal, employer consequences) and speech governance preferences across respondent and context characteristics but also reporting the unconditional shares and averages, respectively.

Finally, we will explore the open responses to the questions on people's definition of hate and their belief about what should and should not be allowed to say on social media quantitatively to potentially identify common themes in definitions and preferences, and compare them with our regular outcome measures.

Other

Copyright © 2011-2023 Center for Open Science (<https://cos.io>) | Terms of Use (https://github.com/CenterForOpenScience/centerforopenscience.org/blob/master/TERMS_OF_USE.md) | Privacy Policy (https://github.com/CenterForOpenScience/centerforopenscience.org/blob/master/PRIVACY_POLICY.md) | Status (<https://status.cos.io/>) | API (<https://developer.osf.io/>)
TOP Guidelines (<http://cos.io/top/>) | Reproducibility Project: Psychology (<https://osf.io/ezcuj/wiki/home/>) | Reproducibility Project: Cancer Biology (<https://osf.io/collections/rpcb/discover>)



(<http://twitter.com/OSFramework>)



(<https://www.facebook.com/CenterForOpenScience/>)



(<https://groups.google.com/forum/#!forum/openscienceframework>)



(<https://www.github.com/centerforopenscience>)



Introduction and informed consent

Welcome!

You are invited to participate in a study conducted by researchers from the Hertie School, one of the world's leading policy schools.

The goal of this survey is to learn about your opinions on various forms of speech online. Please note that we will present various example statements from social media that you may perceive as offensive.

Your participation is voluntary. You are free to not answer any question or to withdraw from the study at any time. This survey should take approximately 15 minutes to complete.

We hope that you would like to take part!

Consent

Before you decide to participate in this study, it is important that you understand why the research is being done and what it will involve. Please take the time to read the following information carefully. Please contact the researchers if you need more information.

Research Project

The title of the research project is "Global Preferences for Hate Speech Moderation". This research project is being conducted by Simon Munzert, Julia Ellingwood, Lisa Oswald and Richard Traunmüller at the Hertie School, Berlin.

Purpose of the Research

This is a study on public opinion and policy preferences in the context of online hate speech. Citizens in multiple countries are surveyed. Your participation is voluntary. Participation involves completion of a survey. You may choose not to answer any or all questions.

Study Procedure

You will be asked a series of questions in an online survey - mostly you will answer by clicking boxes to supply your opinion or experience. You may choose not to answer any questions that you are not comfortable answering. Participation in research is completely voluntary and you can stop any time by simply closing the webpage, at which point data you submitted until then will be recorded. Because the survey is anonymous, we will not be able to determine which survey answers belong to you, so you cannot withdraw information you have provided. You do not waive any legal rights by agreeing to take part in this study.

This survey is anonymous, involves a secure connection, and is being hosted by Qualtrics. Data in transit will always be encrypted and the results of your responses will be stored on an encrypted account which only the researchers will be able to access. However, please note that confidentiality cannot be guaranteed while data are in transit over the Internet.

The survey software provider might use the data for improving product performance and other purposes, excluding personal data. Their services Terms of Service addressing confidentiality may be viewed at <https://www.qualtrics.com/security-statement/>.

Risks and Benefits

Researchers involved in the study will protect your personal information if you accidentally provide any, and others will not be able to connect your responses to personally identifiable information. Any personal information that could identify you will be removed or changed before files are shared with other researchers or results are made public.

The results of this study will be presented at conferences, published in scientific journals, and linked on the researchers' websites. A fully anonymized version of the data will be made public upon publication of the results and will be hosted in perpetuity by the Open Science Framework: <http://osf.io/>. Anyone will be able to view the data from the study once the data is made public, but your answers contain no identifying information and so no one will be able to link any answer to you personally. As such, the survey involves minimal risk to you, and while there are no direct benefits, you will be helping to further scientific understanding of current public opinion.

This project has been reviewed by the Hertie School's Ethics Board for compliance with federal guidelines for research involving human participants. Please feel free to print this consent form for your records.

Contact

Questions about this project may be directed to Simon Munzert at munzert@hertie-school.org. Any concerns can be directed to the Hertie School's dean of research, Kai Wegrich at Dean-of-Research-and-Faculty@hertie-school.org.

I hereby confirm that I am at least 18 years old, that I have read, understood, and agree with the terms detailed above. Do you wish to participate in this survey?

☐ Yes

☐ No

Basic sociodemographics

Please state your gender.

☐ Male

☐ Female

☐ Other

What is your year of birth?


What is the highest level of school you have completed or the highest degree you have received?

- ☐ Some high school or less (Grades 9-11 or Grade 12 with no diploma)
- ☐ High school graduate or GED (includes technical/vocational training that does not count towards college credit)
- ☐ Some college, no degree (includes some community college, associate's degree)
- ☐ Bachelor's degree (e.g., B.A. or B.S.)
- ☐ Some postgraduate or professional schooling, no postgraduate degree
- ☐ Postgraduate or professional degree (includes master's doctorate, medical or law degree)

Which character do you think most closely resembles you?





Which character do you think most closely resembles you?





Which character do you think most closely resembles you?





Are you considering yourself being part of a minority in your country because of your ethnicity or race, religion, gender, sexual orientation, political belief, or another aspect of your identity?

- ☐ Yes, I'm part of a minority in my country.
- ☐ No, I'm not part of a minority in my country.

Do you tend to agree or disagree with the following statements?

	Tend to agree	Tend to disagree	Unsure
We would be better off if everybody were vegan.	☐	☐	☐
The United States of America should be for the American people!	☐	☐	☐
Stay-at-home moms are the best!	☐	☐	☐
It should be fine for men to cry!	☐	☐	☐
It's not my problem if others are hurt by my opinion.	☐	☐	☐
	Tend to agree	Tend to disagree	Unsure
Minority harassment has to disappear from the discourse.	☐	☐	☐
Racism is an overblown issue.	☐	☐	☐
We need strong and independent men!	☐	☐	☐
Women should rule the world!	☐	☐	☐
Immigration makes our society stronger and better!	☐	☐	☐

	Tend to agree	Tend to disagree	Unsure
Climate change is a fundamental threat to humanity.	☐	☐	☐
Our institutions are built on racism.	☐	☐	☐
Climate change is overrated - climate has always changed.	☐	☐	☐
It's my right to eat meat.	☐	☐	☐
	Tend to agree	Tend to disagree	Unsure

Political preferences and behavior

How interested in politics are you?

- ☐ Very interested
- ☐ Moderately interested
- ☐ Slightly interested
- ☐ Not interested at all

Which of these statements comes closest to your opinion?

- ☐ People in this country are **more free** to say what they think than they used to.
- ☐ People in this country are **just as free** to say what they think as they used to.
- ☐ People in this country are **less free** to say what they think than they used to.
- ☐ Don't know

In politics people often talk about the "left" and the "right". On a scale between 1 (furthest left) and 11 (furthest right), where would you place yourself?

1 -
Left

☐

2

☐

3

☐

4

☐

- 5
- ☐
- 6
- ☐
- 7
- ☐
- 8
- ☐
- 9
- ☐
- 10
- ☐
- 11 -
Right
- ☐

Generally speaking, do you usually think of yourself as a

- ☐ Republican
- ☐ Democrat
- ☐ Independent
- ☐ Other
- ☐ None

Do you tend to agree or disagree with the following statements?

	Tend to agree	Tend to disagree	Unsure
Defund the police!	☐	☐	☐
Blue lives matter!	☐	☐	☐
SUVs should be banned!	☐	☐	☐
Democrats are the one and only party.	☐	☐	☐
Religion has done more harm than good to our society.	☐	☐	☐
	Tend to agree	Tend to disagree	Unsure
I am ashamed of my country!	☐	☐	☐
Republicans are the one and only party.	☐	☐	☐

	Tend to agree	Tend to disagree	Unsure
Driving my SUV is pure freedom!	☐	☐	☐
Marriage has to be between a man and a woman.	☐	☐	☐
Gay marriage has to be legal.	☐	☐	☐
	Tend to agree	Tend to disagree	Unsure
A life without faith is an empty life.	☐	☐	☐
The poor are to blame for their poverty.	☐	☐	☐
I am proud of my country!	☐	☐	☐
Tax the rich!	☐	☐	☐
	Tend to agree	Tend to disagree	Unsure

Speech traits

Do you feel more or less free to speak your mind than you used to, or about the same?

- ☐ More free than I used to.
- ☐ Less free than I used to.
- ☐ About the same.
- ☐ Don't know

To what extent do you agree with the following statements?

	Strongly disagree	Somewhat disagree	Somewhat agree	Strongly agree
We must not let people saying unacceptable things be heard. We must act against them.	☐	☐	☐	☐
Sometimes even in a democracy harmful views need to be silenced.	☐	☐	☐	☐

To what extent do you agree with the following statements?

	Strongly disagree	Somewhat disagree	Neither agree nor disagree	Somewhat agree	Strongly agree
I find it easy to take the perspective of others.	☐	☐	☐	☐	☐
I am an empathetic person.	☐	☐	☐	☐	☐
I am good at predicting how others feel.	☐	☐	☐	☐	☐

Online behavior

How often has the following happened to you on the internet, including social media?

	Never	Rarely	Sometimes	Often	Don't know
I have seen views I disagree with.	☐	☐	☐	☐	☐
I have personally been threatened.	☐	☐	☐	☐	☐
I have seen views that make me angry.	☐	☐	☐	☐	☐
I have personally been offended or insulted.	☐	☐	☐	☐	☐
I have witnessed how someone else has been offended, insulted, or threatened.	☐	☐	☐	☐	☐

How often has the following happened to you on the internet, including social media?

	Never	Rarely	Sometimes	Often	Don't know
I posted my opinion on a political issue.	☐	☐	☐	☐	☐
I posted or shared something that I later regretted or felt ashamed of.	☐	☐	☐	☐	☐
I posted or shared something that could be seen as offensive.	☐	☐	☐	☐	☐

Speech governance preferences

Here you can find different arguments people make on current issues.

Which side do you agree with more?

People should be able to speak their minds freely online.

☐

People should be able to feel welcome and safe online.

☐

Which side do you agree with more?

Online services should not be responsible for the content users post on their site, even when it's harassing.

☐

Online services have a responsibility to step in when harassing behavior occurs on their site.



Which side do you agree with more?

Government should prevent people from engaging in hate speech against certain groups in public.



People should be allowed to express unpopular opinions in public, even those that are deeply offensive to other people.



To what extent, if at all, do you think each of the following groups have responsibility to act against online hate speech?

	No responsibility at all	Rather no responsibility	Some responsibility	Full responsibility	Don't know
People who are victims of online hate speech	☐	☐	☐	☐	☐
Other users who witness the behavior	☐	☐	☐	☐	☐
Online services such as social media platforms	☐	☐	☐	☐	☐
Lawmakers	☐	☐	☐	☐	☐
Justice system	☐	☐	☐	☐	☐
Employers of writers of hate speech	☐	☐	☐	☐	☐

Please indicate who - if any - should mainly be responsible for regulating content on social media platforms:

- ☐ The social media companies
- ☐ The national government
- ☐ Both the social media companies and the government
- ☐ There should be no corporate or government regulation of content on social media

Vignettes introduction



You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content.

How would you act if you had to decide what happens with the messages?

Please read the statements carefully and answer a few questions for each of them. There is no right or wrong, please answer spontaneously.

Important note: We also want to point out that the contents of some of these messages may be unpleasant or repugnant to you. If you do not want to see any more such messages, you can skip these questions without answering.

- ☐ I have read and understood those instructions.



You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content.

Imagine you were a content moderator for a social platform that vows to promote free speech and takes a critical stance against any form of censorship. How would you act if you had to decide what happens with the messages?

Please read the statements carefully and answer a few questions for each of them. There is no right or wrong, please answer spontaneously.

Important note: We also want to point out that the contents of some of these messages may be unpleasant or repugnant to you. If you do not want to see any more such messages, you can skip these questions without answering.

☐ I have read and understood those instructions.



You will now see different examples of conversations that could typically be found on social media. Many of these conversations contain controversial language and content.

Imagine you were a content moderator for a social platform that vows to protect their users from hate speech and takes a committed stance against hateful content on their pages. How would you act if you had to decide what happens with the messages?

Please read the statements carefully and answer a few questions for each of them. There is no right or wrong, please answer spontaneously.

Important note: We also want to point out that the contents of some of these messages may be unpleasant or repugnant to you. If you do not want to see any more such messages, you can skip these questions without answering.

- ☐ I have read and understood those instructions.

Vignettes

Conversation 1/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the **SECOND** post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech

- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 2/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the **SECOND post**.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed

- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 3/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the SECOND post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 4/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the SECOND post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 5/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the **SECOND post**.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 6/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the SECOND post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 7/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the SECOND post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Conversation 8/8

Please consider the entire conversation before you answer the questions.



After having read the entire conversation, consider the SECOND post.

Do you think the **second post** constitutes hate speech, offensive speech, or neither of them?

- ☐ Hate speech
- ☐ Offensive but not hate speech
- ☐ Neither offensive nor hate speech

Do you think the **second post** should be removed from the platform or not?

- ☐ Post should be removed
- ☐ Post should not be removed

Do you think the **sender of the second post** should be banned from the platform or not?

- ☐ Person should be banned
- ☐ Person should not be banned

Do you think that the **sender of the second post** should face legal consequences, such as a fine?

- ☐ Person should face legal consequences
- ☐ Person should not face legal consequences

Do you think that the **sender of the second post** should face consequences from their employer, such as being fired?

- ☐ Person should face consequences by employer
- ☐ Person should not face consequences by employer

Post experiment

In the previous questions we asked you to evaluate various conversations as they could have happened on a social platform. When answering those messages, which kind of social platform did you have in mind?

- ☐ A social platform that vows to promote free speech and takes a critical stance against any form of censorship
- ☐ A social platform that vows to protect their users from hate speech and takes a committed stance against hateful content on their pages
- ☐ Another kind of social platform
- ☐ I did not have a particular kind of social platform in mind

End

Before we come to the end, we have two last questions.

People have different ideas about what constitutes "hate speech." What about you - how would you personally define hate speech?



And, in your own words, what do you think:

What - if anything - should not be allowed to say on social media?



Thank you very much for participating in the survey.

With this study, we aim to understand what people think about speech online and how they prefer hate speech to be moderated and regulated. We believe that our research is important to inform the scientific and public debate about everybody's opinions about the issue. In particular, we are interested in how social media users think about the dilemma between protecting free speech online, and protecting people from potential harm from speech online.

The social media messages you saw in the survey were composed of different features, such as sender and target characteristics and messages of different type and degree. The statements were artificially created by the research team and as such were not based on real users or content from online platforms. The content you saw can be considered toxic and offensive, and some of the statements, if they were expressed online, might violate community standards on social media or result in criminal prosecution. **Neither the researchers nor the Hertie School endorse the views depicted in the messages and pictures, and we kindly ask you not to re-use the material online.** It was necessary, however, to present you with statements of this kind to be able to inform our research question.

Furthermore, in the survey you were randomly assigned to one of three experimental treatments, which introduced the task to evaluate the social media messages in different ways. In one version of the instructions, we asked participants to imagine they were a content moderator for a social platform that vows to promote free speech and takes a critical stance against any form of censorship. In another version of the instructions, we asked participants to imagine they were a content moderator for a social platform that vows to protect their users from hate speech and takes a committed stance against hateful content on their pages. In the third version, none of those frames were presented. We used this experimental setup to study whether participants' perceptions and evaluations of hate speech can be affected by the information provided in the instructions.

If you have any questions regarding the survey, please contact the principal investigator of the study, Prof. Simon Munzert, Hertie School: munzert@hertie-school.org.

Thank you!

Powered by Qualtrics